

Data Science: Capstone

PH125.9x

HarvardX Professional Certificate in Data Science

December 2020

Student: Miguel Angel Bustos Sáez

Instructor: Rafael Irizarry

Introduction

This inform it's a recommendation system, based in Random Forest¹ for categorical data and RMSE² for numerical data, using Edx dataset joined it with Imdb ratings it gives important information to get a better forecast approximation. The steps of the project are:

1. R Libraries

- Tidyverse
- Magrittr
- Lubridate
- randomForest
- caret
- readr

2. Methods and analysis

3. Datasets structures

4. Data model

5. Model results

- Random forest accuracy (For categorical variable)
- RMSE forecast accuracy (For numerical variable)

6. Visualization

7. Conclusion

¹https://en.wikipedia.org/wiki/Root_mean_square

²https://en.wikipedia.org/wiki/Random_forest

2 Methods and analysis

Data cleaning

The data cleaning process was dividing the Edx dataset in movies.csv and ratings.csv and doing an inner_joining with imdb dataset, that provided more information to the model; votes numbers, imdb average and Id user. Edx dataset had a separate **I** symbol and it had all genres in just one column, so the separate **I** symbol was removed using the separate function, and str_c function ordered each genre in each column:

genero_1	genero_2	genero_3	genero_4	genero_5	genero_6	genero_7	genero_8	genero_9	genero_10
Action	Adventure	Comedy	Crime	Drama	Film-Noir	Horror	Mystery	Thriller	Western
Action	Adventure	Comedy	Crime	Drama	Film-Noir	Horror	Mystery	Thriller	Western
Action	Adventure	Comedy	Crime	Drama	Film-Noir	Horror	Mystery	Thriller	Western
Action	Adventure	Comedy	Crime	Drama	Film-Noir	Horror	Mystery	Thriller	Western
Action	Adventure	Comedy	Crime	Drama	Film-Noir	Horror	Mystery	Thriller	Western
Action	Adventure	Comedy	Crime	Drama	Film-Noir	Horror	Mystery	Thriller	Western
Action	Adventure	Comedy	Crime	Drama	Film-Noir	Horror	Mystery	Thriller	Western
Action	Adventure	Comedy	Crime	Drama	Film-Noir	Horror	Mystery	Thriller	Western
Action	Adventure	Comedy	Crime	Drama	Film-Noir	Horror	Mystery	Thriller	Western
Action	Adventure	Comedy	Crime	Drama	Film-Noir	Horror	Mystery	Thriller	Western
Action	Adventure	Comedy	Crime	Drama	Film-Noir	Horror	Mystery	Thriller	Western
Action	Adventure	Comedy	Crime	Drama	Film-Noir	Horror	Mystery	Thriller	Western
Action	Adventure	Comedy	Crime	Drama	Film-Noir	Horror	Mystery	Thriller	Western
Action	Adventure	Comedy	Crime	Drama	Film-Noir	Horror	Mystery	Thriller	Western
Action	Adventure	Comedy	Crime	Drama	Film-Noir	Horror	Mystery	Thriller	Western
Action	Adventure	Comedy	Crime	Drama	Film-Noir	Horror	Mystery	Thriller	Western

To fill some empty cells, in year variable, it processes to take the average year, and fill some cells with the average year 1992. The genres are ordered and the missing pieces are filled with `NA`:

Titulo	Genero	genero_1	genero_2	genero_3	genero_4	genero_5	genero_6	genero_7	genero_8	genero_9	genero_10
All	All	All	All	All	All	All	All	All	All	All	All
Amelie (Fabuleux de...	Comedy Romance	Comedy	Romance	NA	NA	NA	NA	NA	NA	NA	NA
Wild Strawberries (S...	Drama	Drama	NA	NA	NA	NA	NA	NA	NA	NA	NA
Teddy Bear (Mis) (19...	Comedy Crime	Comedy	Crime	NA	NA	NA	NA	NA	NA	NA	NA
Talk to Her (Hable c...	Drama Romance	Drama	Romance	NA	NA	NA	NA	NA	NA	NA	NA
Lord of the Rings: T...	Adventure Fantasy	Adventure	Fantasy	NA	NA	NA	NA	NA	NA	NA	NA
City of God (Cidade ...	Action Adventure...	Action	Adventure	Crime	Drama	Thriller	NA	NA	NA	NA	NA
Spanish Apartment, ...	Comedy Drama R...	Comedy	Drama	Romance	NA	NA	NA	NA	NA	NA	NA
Finding Nemo (2003)	Adventure Anima...	Adventure	Animation	Children	Comedy	NA	NA	NA	NA	NA	NA
Lost in Translation (...)	Comedy Drama R...	Comedy	Drama	Romance	NA	NA	NA	NA	NA	NA	NA
M. Hulot's Holiday (...)	Comedy	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
Strada, La (1954)	Drama	Drama	NA	NA	NA	NA	NA	NA	NA	NA	NA

3 Datasets structures

It is edx dataset, where genres variable was separated by this **I** symbol:

```
> str(edx)
'data.frame':   9000060 obs. of  6 variables:
 $ userId      : int  1 1 1 1 1 1 1 1 1 1 ...
 $ movieId     : num  122 185 292 316 329 355 362 364 370 377 ...
 $ rating      : num  5 5 5 5 5 5 5 5 5 5 ...
 $ timestamp   : int  838985046 838983525 838983421 838983392 838983392 83898
 $ title       : chr   "Boomerang (1992)" "Net, The (1995)" "Outbreak (1995)"
 $ genres      : chr   "Comedy|Romance" "Action|Crime|Thriller" "Action|Drama|
```

And now this is our data, an inner_join combination between Edx with Imdb datasets, the recommendation systems are based on this data:

```
> str(data)
'data.frame':   25475592 obs. of  9 variables:
 $ Id_usuario   : int  1 1 1 1 1 1 1 1 1 1 ...
 $ Id_pelicula  : num  296 306 307 665 899 ...
 $ Indice       : num  5 3.5 5 5 3.5 4 3.5 3.5 5 4 ...
 $ Marca_temporal: int  1147880044 1147868817 1147868828 1147878820
 $ Titulo       : chr   "Pulp Fiction (1994)" "Three Colors: Red (Tr
 "Underground (1995)" ...
 $ Genero       : chr   "Comedy|Crime|Drama|Thriller" "Drama" "Drama
 $ Ano         : num  1994 1994 1993 1995 1952 ...
 $ Promedio_Idbm : num  8.9 8.1 7.9 8.1 8.3 7 7.6 8.2 8.2 8.1 ...
 $ Numero_Votos : num  1785576 88165 86831 54330 213897 ...
```

It is vital to have the most information as possible in genres, having well ordered in different columns, for that objective, this data is a combination between Edx and Imdb datasets, to obtain a major accuracy in the predictions.

4 Data model

The data, was saved as `data.rda`, taken the major and equal 4 índice or number rating. The *training_u* is a length of 50% and *testing_u* vectors with the 50% of users. *Training* it's a filter of the `id_users` of *training_u*, selecting the variables, and adding it a `generos` column, here it is the code.

The data has been divided in two objects; *training_u* and *testing_u*:

```
data = data %>% filter(Genero != "(no genres listed)" & Indice >= 4)

Usuarios = unique(data$Id_usuario)
training_u = sample(Usuarios, length(Usuarios)*0.50,replace = F)
testing_u = Usuarios[-training_u]

training = data %>%
  filter(Id_usuario %in% training_u)%>%
  select(Indice,Ano,Generos,Marca_temporal,Promedio_Idbm,Numero_Votos)%>%
  mutate(Generos = factor(Generos))
```

And here it is a random forest and the `set.seed`:

```
set.seed(0)
modelo = randomForest(Generos ~ ., data = training,
                      ntree=30,
                      method="class",
                      norm.votes=FALSE,
                      do.trace=10,
                      proximity = FALSE,
                      importance=TRUE)

muestra_2 = list(testing_u, modelo)
save(muestra_2, file="muestra_2.rda")
```

5 Model results

I choose two algorithms approximation; Random Forest for categorical data, and RMSE for numerical data.

5.1 Random forest for categorical data

This code with the OOB estimate error rate of 68.29%:

```
[[2]]

Call:
randomForest(formula = Generos ~ ., data = training, ntree = 30, method = "class", norm.votes = FALSE,
y = FALSE, importance = TRUE)
Type of random forest: classification
Number of trees: 30
No. of variables tried at each split: 2

OOB estimate of error rate: 68.29%
Confusion matrix:
```

	Action	Adventure	Animation	Children	Comedy	Crime	Documentary	Drama	Fantasy	Film-Noir	Horror
Action	512370	228677	13410	7724	94154	98728	286	161530	47469	461	22819
Adventure	252618	375029	53725	65513	113541	10361	883	119752	114518	614	6330
Animation	17751	71411	82494	79021	59403	1858	409	26315	37051	6	745
Children	12503	81347	75452	107744	76552	1533	183	35131	49319	1	191
Comedy	100586	128853	51290	70650	827068	86070	4689	298922	70374	441	21265
Crime	145489	15476	3073	2666	115719	319131	597	274911	3194	16138	22241
Documentary	348	834	432	182	6116	482	74958	5411	45	1	112
Drama	172623	95054	18948	39400	263126	225139	3207	1447829	49027	7781	24281
Fantasy	50863	121424	42073	58388	88629	3128	79	63176	161122	191	9811
Film-Noir	2138	2064	11	0	480	20051	0	13314	160	19279	198
Horror	30518	6147	977	210	26557	18290	102	37531	11261	404	125570
IMAX	49995	39551	7867	9021	4185	13046	770	28279	17496	0	1253
Musical	1893	15913	21880	32417	43240	468	2974	32790	23057	7	2043
Mystery	35349	11316	2033	1744	13538	62818	218	106683	8215	13390	20616
Romance	35371	30897	13498	11492	266865	12005	92	341426	42826	2115	2536
Sci-Fi	205633	144221	13290	9905	46111	12783	197	92452	14009	596	36395
Thriller	264094	71076	1693	210	46188	185235	49	274035	12680	14218	68646
War	48808	14358	1097	280	30158	2972	2219	151361	1110	4	172
Western	23541	14382	262	427	10816	1458	11	25257	170	0	516

	IMAX	Musical	Mystery	Romance	Sci-Fi	Thriller	War	Western	class.error
Action	36579	2046	20263	34381	185093	221049	40743	12809	0.7056345
Adventure	36061	11481	7593	22291	133220	77956	10876	13922	0.7370587
Animation	4761	19777	1661	13845	10700	1309	1547	174	0.8082596
Children	4841	21371	1501	12381	9572	294	256	313	0.7803317
Comedy	2558	29193	9130	201026	44708	45872	28139	8915	0.5925270
Crime	8136	703	38647	9632	9930	184982	3075	897	0.7283152
Documentary	645	2439	190	49	119	39	2156	7	0.2073389
Drama	12061	16962	67824	216055	72354	201685	107479	17568	0.5266062
Fantasy	18255	12641	6410	32083	9886	12427	1853	168	0.7673688
Film-Noir	0	11	9915	2955	346	12815	10	4	0.7698057
Horror	1798	1586	24763	3311	33381	63098	102	405	0.6746984
IMAX	51795	2986	1812	2749	29390	12915	1743	43	0.8115833
Musical	2153	43066	64	21197	1351	310	518	191	0.8246013
Mystery	3927	93	131863	11906	24635	107965	517	553	0.7634231
Romance	2843	19404	13248	271696	15874	24868	20140	3285	0.7596634
Sci-Fi	30308	1249	23197	10969	266407	114334	1184	2633	0.7403119
Thriller	13786	607	118491	27622	109244	436276	8840	1444	0.7362989
War	1887	605	473	16149	1396	10975	102123	1778	0.7367455
Western	68	258	788	5808	1025	1881	1257	38225	0.6969877

Errors have different class.errors fluctuations: 0.70 for action, 0.73 for adventure, 0.80 for animation, etc.

Unique() function shows the id_users, the 1990 user has been chosen here:

```
unique(testing$Id_usuario)

testing_1900 = testing %>% filter(Id_usuario == 1900) %>%
  select(Indice, Ano, Genero, Marca_temporal, Promedio_Idbm, Numero_Votos) %>%
  mutate(Genero = factor(Genero))
testing_1900$Genero = fct_expand(testing_1900$Genero, levels(factor(testing$Genero))[2:20])
```

Genres and movies suggestions for the 1990 user

	Indice	Ano	Genero	Marca_temporal	Promedio_Idbm	Numero_Votos
1	3.0	1995	Comedy Romance	1301834435	6.3	36058
2	2.0	1995	Action Romance Western	1301834430	7.2	169053
3	2.5	1995	Horror Sci-Fi	1301834464	5.8	72994
4	3.5	1977	Action Adventure Sci-Fi	1301834854	8.6	1207095
5	4.5	1994	Comedy Crime Drama Thriller	1301835024	8.9	1785576
6	4.5	1994	Crime Drama	1301835077	9.3	2285846
7	5.0	1994	Comedy Drama Romance War	1301834954	8.8	1761399
8	3.5	1994	Action Comedy	1301834425	6.5	95555
9	3.0	1993	Thriller	1301835195	7.8	262807
10	4.0	1991	Action Sci-Fi	1301834882	8.5	977571
11	4.0	1972	Crime Drama	1301835761	9.2	1577887
12	3.0	1941	Film-Noir Mystery	1301834493	8.0	147039
13	3.5	1964	Comedy Drama Musical Romance	1301834532	7.8	85000
14	3.0	1988	Action Crime Thriller	1301835145	8.2	769718
15	2.5	1996	Comedy Drama	1301834537	7.2	74180
16	4.0	1980	Action Adventure Sci-Fi	1301834826	8.7	1135176
17	4.0	1962	Adventure Drama War	1301834486	8.3	262358

Showing 1 to 18 of 86 entries, 6 total columns

Genres suggestions for the 1990 user:

```
> generos_sugerir = names(sort(table(prediccion_1900),decreasing = T)[1:3])
> generos_sugerir
[1] "Drama"      "Adventure" "Comedy"
```

Movies suggestions for the 1990 user:

```
i1 = data %>% filter(Genero %in% "Drama" & Promedio_Idbm >= 8 & Ano >= 2015) %>% distinct(Titulo)
```

	Titulo
1	Room (2015)
2	Lion (2016)
3	Dangal (2016)
4	Bohemian Rhapsody (2018)
5	Wonder (2017)
6	Shoplifters (2018)
7	Fly Away Solo (2015)
8	The Wild Pear Tree (2018)

5.2 RMSE for numerical data

Here it is a RMSE approximation, it's 28.97 the numerical approximation to obtain a recommendation movie, this approximation is obtained based on a confusion matrix³ and RMSE result:

```
73 unique(testing$Id_usuario)
74 testing_1900 = testing %>% filter(Id_usuario == 1900) %>% select(Indice,Ano,Gen
75 testing_1900$Genero = fct_expand(testing_1900$Genero,levels(factor(testing$Gene
76
77 set.seed(100)
78 prediccion_1900 = predict(modelo, newdata = testing_1900, type = "class") # typ
79
80 confusionMatrix(prediccion_1900, testing_1900$Generos)
81
82 cm = table(prediccion_1900,testing_1900$Generos)
83 i =
84   j = 1:dim(cm)[2]
85
86 se = 0;
87 for (i in 1:dim(cm)[1])
88 {
89   for (j in 1:dim(cm)[2])
90   {
91     se = se + cm[i,j] * (i-j)^2;
92   }
93 }
94
95 mse = se / sum(sum(cm))
96 (rmse = sqrt(mse))
97
98
99
71:35 Forecast category <
Console Terminal x
~/Desktop/Escritorio/A Data/2 THIS/Recomendacion/Miguel/Migu/ ↵
+ {
+   se = se + cm[i,j] * (i-j)^2;
+ }
+ }
> mse = se / sum(sum(cm))
> (rmse = sqrt(mse))
[1] 28.97011
```


Continuous with our testing, here we can see again our 1990 user, we obtain in the descendant arrange; 7.87 for drama and 6.4 for comedy genres:

```
> testing_1900 %>%
+   filter(Generos %in% generos_sugerir) %>%
+   group_by(Generos) %>%
+   summarise(idbm = mean(Promedio_Idbm)) %>%
+   arrange(desc(idbm))
# A tibble: 2 x 2
  Generos idbm
<fct>    <dbl>
1 Drama   7.87
2 Comedy  6.4
>
```

³https://en.wikipedia.org/wiki/Confusion_matrix

6 Visualization

In this histogram we have the average ratings from 2000 to 2019 of data dataset:




*Ratings at moving, are also knowing as **content rating**, it rates the suitability on TV broadcasts, movies, internet, music, comics books or video games to its audience ⁴.*

Is very important have the information of content rating because provide guidance to consumers, particularly parents, to help them decide whether or not to watch any movie. For example, Netflix decide maturity ratings by country and for kids 7+, for teens 13+ and adult 16+, 18+ . The good of Netflix is that the user can choose maturity ratings or block shows⁵. In Netflix the average of watching tv is two hours per user, concluding all, data and visualization is very important to take decisions.

⁴[https://en.wikipedia.org/wiki/Content_rating#:~:text=A%20content%20rating%20\(also%20known,suitable%20to%20view%20said%20media](https://en.wikipedia.org/wiki/Content_rating#:~:text=A%20content%20rating%20(also%20known,suitable%20to%20view%20said%20media)


⁵<https://help.netflix.com/en/node/2064>

17 genres rated 1 to 5 (From 2015 to 2019):



Here are some genres, in the years 2015 to 2019. Action, comedy and drama rocks!

Horror is very fluctuated, mystery is highly rated, romance is very scattered:



Conclusion

The main process in this project is the movies.csv, many movies were repeated, and the process was fixing this database without movies repeated, obtained 62.423 observations and three variables. In the column genres it continuous with more than one genre, but in the code, this | character was removed. Also, this project has an IMDB ratings data set 1.076.066 observations and three variables; ID, average rating and number of votes. Both datasets were mutate in the code.

The courses helped me to have tools to sort the data, unite and based on this establish the RMSE and the Random Forrest to obtain effective predictions, and for this, I decided to establish these two algorithms, both categorical and numerical, thus having a more global clarity about of the data, with this, the histograms and boxplots presented in a more clear and efficient visualization.

The edx team, always where very efficient in all responses, Rafael Irizarry classes where very nutritive in any knowledge and great tips to use the R function in different examples. I conclude this course was very hard, it takes me a lot of hours to understand some examples, I spent almost two years to finish this whole program, and I will continuous studying statistics and R language.

I am very grateful for this enriching course and helpful team.

Thanks!

Miguel Angel Bustos Sáez

